

# Practical Machine Learning-Week4 Assignment

Ahad Zaman Khan

2022-09-26

## Loading and processing The Data

I downloaded the data and reading it from local drive.

```
library(tidyverse)
list.files()

## [1] "Assignment.R"          "pml-testing.csv"      "pml-training.csv"
## [4] "Rplot.png"            "week4_assignment.html" "week4_assignment.Rmd"
## [7] "week4_assignment_cache" "week4_assignment_files"

pml_train <- read_csv("pml-training.csv")
```

First column of the data is row numbers. Lets remove it from the data. I also check how the class variable is distributed.

```
pml_train <- pml_train[ , -1]

# names(pml_train)

length(unique(pml_train$classe))

## [1] 5

unique(pml_train$classe)

## [1] "A" "B" "C" "D" "E"

# find distribution of class variable-
pml_train %>% group_by(classe) %>%
  summarise( no_obs = n(), prcnt = 100* n()/nrow(.) )
```

```
## # A tibble: 5 x 3
##   classe no_obs prcnt
##   <chr>   <int> <dbl>
## 1 A       5580  28.4
## 2 B       3797  19.4
## 3 C       3422  17.4
## 4 D       3216  16.4
## 5 E       3607  18.4
```

Lets check for missing values in the data-

```
# function to check a summary for "NA"s
na_count <-
  function(z) {
    num_obs = nrow(z)
    y = map_df(z, ~ sum(is.na(.x)))
```

```

y = pivot_longer(y,
                  everything(),
                  names_to = "variable",
                  values_to = "num_na") %>%
  filter(num_na > 0)
y$prcnt = y$num_na *100 /num_obs

return(y)
}

na_count(pml_train) %>% arrange(-prcnt)

```

```

## # A tibble: 100 x 3
##   variable          num_na prcnt
##   <chr>            <int> <dbl>
## 1 skewness_roll_belt    19225  98.0
## 2 skewness_roll_dumbbell 19220  98.0
## 3 skewness_pitch_dumbbell 19217  97.9
## 4 kurtosis_roll_belt    19216  97.9
## 5 kurtosis_picth_belt   19216  97.9
## 6 kurtosis_yaw_belt     19216  97.9
## 7 skewness_roll_belt.1  19216  97.9
## 8 skewness_yaw_belt     19216  97.9
## 9 max_roll_belt        19216  97.9
## 10 max_picth_belt      19216  97.9
## # ... with 90 more rows

```

100 variables out of 159 variable has 97.93% to 97.97% records are missing "NA".I am omitting this variables and creating a new data named pml\_train2. lets find out which variables does not have missing values:

```

pml_train2 <- pml_train %>%
  select( -na_count(pml_train)$variable )

```

lets find the variables that are not numeric:

```

pml_train2 %>% select_if(~!is.numeric(.)) %>% names()

```

```
## [1] "user_name"      "cvtd_timestamp" "new_window"     "classe"
```

4 columns are not numerical. "classe" is dependent / predict variable. Ignoring other 3 variables.

```

pml_train2 <- pml_train2 %>%
  select( - c("user_name", "cvtd_timestamp", "new_window") )

```

## Partitioning the data

```

library(caret)

train_sample <- createDataPartition(y = pml_train2$classe, p = 0.7, list = FALSE)

training <- pml_train2[train_sample,]
testing <- pml_train2[-train_sample,]

```

## Prediction models:

### DECISION TREE

```
DT_model <- train(classe~.,
                  data = training, method = "rpart")

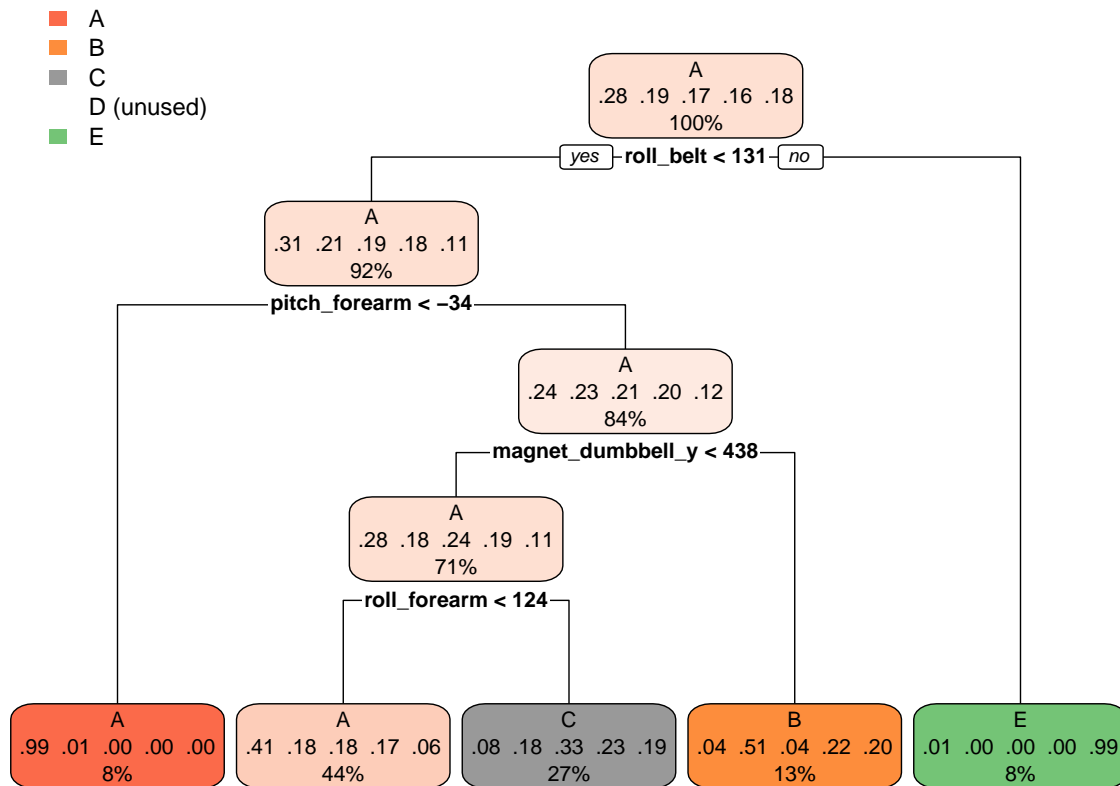
DT_prediction <- predict(DT_model, testing)

confusionMatrix(DT_prediction, as.factor(testing$classe))

## Confusion Matrix and Statistics
##
##           Reference
## Prediction   A    B    C    D    E
##           A 1528  468  491  430  157
##           B   27  400   35  188  147
##           C  116  271  500  346  260
##           D    0    0    0    0    0
##           E    3    0    0    0  518
##
## Overall Statistics
##
##               Accuracy : 0.5006
##               95% CI : (0.4877, 0.5135)
##      No Information Rate : 0.2845
##      P-Value [Acc > NIR] : < 2.2e-16
##
##               Kappa : 0.3469
##
##  McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##               Class: A Class: B Class: C Class: D Class: E
## Sensitivity           0.9128 0.35119 0.48733 0.0000 0.47874
## Specificity           0.6329 0.91635 0.79564 1.0000 0.99938
## Pos Pred Value        0.4971 0.50188 0.33490      NaN 0.99424
## Neg Pred Value        0.9481 0.85476 0.88024 0.8362 0.89485
## Prevalence            0.2845 0.19354 0.17434 0.1638 0.18386
## Detection Rate        0.2596 0.06797 0.08496 0.0000 0.08802
## Detection Prevalence  0.5223 0.13543 0.25370 0.0000 0.08853
## Balanced Accuracy      0.7728 0.63377 0.64148 0.5000 0.73906
```

We can see prediction accuracy is only 52.05%.

```
rpart.plot:: rpart.plot(DT_model$finalModel,
                        roundint = FALSE)
```



## RANDOM FOREST MDOEL

```
rm_md1 <- train(classe~., data = training, method = "rf", ntree = 250)
```

```
rf_predict <- predict(rm_md1, testing)
confusionMatrix(rf_predict, factor(testing$classe))
```

## Confusion Matrix and Statistics

##

## Reference

Prediction	A	B	C	D	E
A	1674	2	0	0	0
B	0	1137	2	0	0
C	0	0	1024	1	0
D	0	0	0	962	0
E	0	0	0	1	1082

##

## Overall Statistics

##

## Accuracy : 0.999

## 95% CI : (0.9978, 0.9996)

## No Information Rate : 0.2845

## P-Value [Acc > NIR] : < 2.2e-16

##

## Kappa : 0.9987

##

```
## McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      1.0000   0.9982   0.9981   0.9979   1.0000
## Specificity      0.9995   0.9996   0.9998   1.0000   0.9998
## Pos Pred Value   0.9988   0.9982   0.9990   1.0000   0.9991
## Neg Pred Value   1.0000   0.9996   0.9996   0.9996   1.0000
## Prevalence       0.2845   0.1935   0.1743   0.1638   0.1839
## Detection Rate   0.2845   0.1932   0.1740   0.1635   0.1839
## Detection Prevalence 0.2848   0.1935   0.1742   0.1635   0.1840
## Balanced Accuracy 0.9998   0.9989   0.9989   0.9990   0.9999
```

### Gradient boosting model:

```
gbm_model <- train(classe~., data = training, method = "gbm", verbose = FALSE )
gbm_model$finalModel
```

```
## A gradient boosted model with multinomial loss function.
## 150 iterations were performed.
## There were 55 predictors of which 55 had non-zero influence.
```

```
gbm_predict <- predict(gbm_model, testing)
confusionMatrix(gbm_predict, factor(testing$classe))
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction    A    B    C    D    E
##           A 1674    2    0    0    0
##           B    0 1135    3    0    0
##           C    0    2 1020    3    0
##           D    0    0    3  957    1
##           E    0    0    0    4 1081
##
## Overall Statistics
##
##           Accuracy : 0.9969
##           95% CI : (0.9952, 0.9982)
##           No Information Rate : 0.2845
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9961
##
## McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      1.0000   0.9965   0.9942   0.9927   0.9991
## Specificity      0.9995   0.9994   0.9990   0.9992   0.9992
## Pos Pred Value   0.9988   0.9974   0.9951   0.9958   0.9963
```

## Neg Pred Value	1.0000	0.9992	0.9988	0.9986	0.9998
## Prevalence	0.2845	0.1935	0.1743	0.1638	0.1839
## Detection Rate	0.2845	0.1929	0.1733	0.1626	0.1837
## Detection Prevalence	0.2848	0.1934	0.1742	0.1633	0.1844
## Balanced Accuracy	0.9998	0.9979	0.9966	0.9960	0.9991

I can see both Random forest and Gradient boosting model are highly accurate. However, Random forest is better performing and I will use this as my prediction model.

## LOADING TESTING DATA and cleaning in same manner as test data

```
pml_test <- read_csv("pml-testing.csv")

pml_test <- pml_test[, -1]

pml_test <- pml_test %>%
  select( -na_count(pml_train)$variable )

pml_test <- pml_test %>%
  select( - c("user_name", "cvtd_timestamp", "new_window") )
```

## Final Prediction:

```
rf_final_predition <- predict(rm_mdl, pml_test)

rf_final_predition

## [1] B A B A A E D B A A B C B A E E A B B B
## Levels: A B C D E
```